

# Data Structures and Algorithms

## Lecture 12: **B<sup>+</sup> Trees I**

Department of Computer Science & Technology  
United International College

# Agenda

- Why B<sup>+</sup> Trees?
- B<sup>+</sup> Tree Introduction
- B<sup>+</sup> Tree Operations
  - Search
  - Insertion
  - Deletion

# Motivation

- An AVL tree with  $N$  nodes is an excellent data structure for searching, indexing, etc.
  - The Big-Oh analysis shows most operations finish within  $O(\log(N))$  time
- The theoretical conclusion works as long as the entire structure can fit into the main memory
- When the data size is too large and has to reside on disk, the performance of AVL trees may deteriorate rapidly

# A Practical Example

- A 500-MIPS machine, with 7200 RPM hard disk
  - 500 million instruction executions, and approximately 120 disk accesses each second
- The machine is shared by 20 users
  - Thus for each user, can handle  $120/20=6$  disk access/sec
- A database with 10,000,000 items,
  - 256 bytes/item (assume it doesn't fit in main memory)
  - The typical searching time for one user
    - A successful search need  $\log_2(10,000,000) = 24$  disk access,
    - Takes around  $24/6=4$  sec.
    - This is way too slow!!
- We want to reduce the number of disk access to a very small constant



# From Binary to M-ary

- Idea: allow a node in a tree to have many children
  - Less disk access = smaller tree height = more branching
- As branching increases, the depth decreases
- An M-ary tree allows M-way branching
  - Each internal node has at most M children
- A complete M-ary tree has height that is roughly  $\log_M(N)$  instead of  $\log_2(N)$ 
  - if  $M = 20$  then  $\log_{20}(2^{20}) < 5$
  - Thus, we can speedup the search significantly

# M-ary Search Tree

- Binary search tree has one key to decide which of the two branches to take
- M-ary search tree needs  $M-1$  keys to decide which branch to take
- M-ary search tree should be balanced in some way too
  - We don't want an M-ary search tree to degenerate to a linked list, or even a binary search tree
    - Thus, require that each node is at least **half** full!

# B<sup>+</sup> Tree

- A B<sup>+</sup>-tree of order M ( $M > 3$ ) is an M-ary tree with the following properties:
  1. The data items are stored in leaves
  2. The root is either a leaf or has between two and M children
  3. The non-leaf nodes store up to M-1 keys to guide the searching; key  $i$  represents the smallest key in subtree  $i+1$
  4. All non-leaf nodes (except the root) have between  $\lceil M/2 \rceil$  and M children
  5. All leaves are at the same depth and have between  $\lceil L/2 \rceil$  and L data items, for some L (usually  $L \ll M$ , but we will assume  $M=L$  in most examples)

There are various definitions of B-trees that differ in minor ways.  
The above definition is one of the popular forms.

# Keys in Internal Nodes

- Which keys are stored at the internal nodes?
  - There are several ways to do it. Different books adopt different conventions.
- We will adopt the following convention:
  - key  $i$  in an internal node is the smallest key in its  $i+1$  subtree (i.e. right subtree of key  $i$ )
- Even following this convention, there is no unique B<sup>+</sup>-tree for the same set of records.



# B<sup>+</sup> Tree Example 1 (M=L=5)

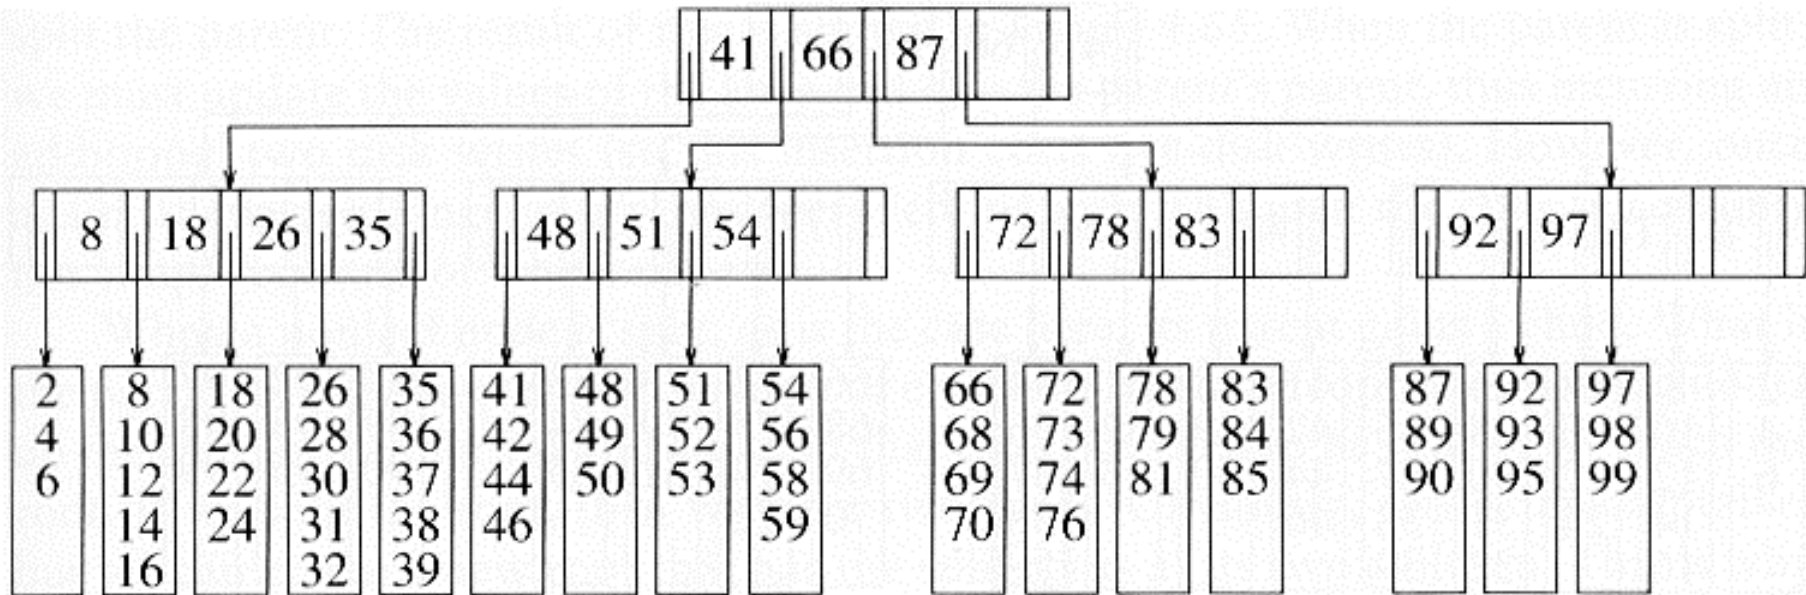



Figure 4.62 B-tree of order 5

- Records are stored at the leaves (we only show the keys here)
- Since  $L=5$ , each leaf has between 3 and 5 data items
- Since  $M=5$ , each nonleaf nodes has between 3 to 5 children
- Requiring nodes to be **half full** guarantees that the B+ tree does not degenerate into a simple binary tree

# B<sup>+</sup> Tree Example 2 (M=4,L=3)



- We can still talk about **left** and **right child pointers**
- E.g. the left child pointer of N is the same as the right child pointer of J
- We can also talk about the **left subtree** and **right subtree** of a key in internal nodes

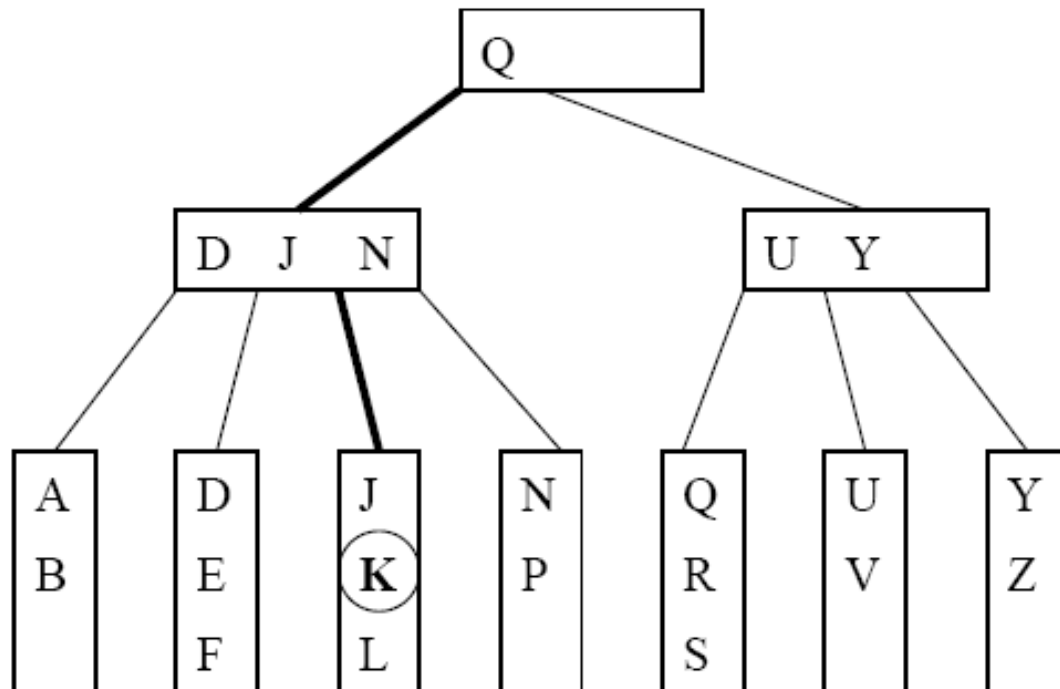
# B+ Tree in Practical Usage



- Each internal node/leaf is designed to fit into one I/O block of data. An I/O block usually can hold quite a lot of data. Hence, an internal node can keep a lot of keys, i.e., large  $M$ . This implies that the tree has **only a few levels** and **only a few disk accesses** can accomplish a search, insertion, or deletion.
- B<sup>+</sup>-tree is a popular structure used in commercial databases. To further speed up the search, **the first one or two levels of the B<sup>+</sup>-tree are usually kept in main memory**.
- The disadvantage of B<sup>+</sup>-tree is that most nodes will have less than  $M-1$  keys most of the time. This could lead to severe **space wastage**. Thus, it is not a good dictionary structure for data in main memory.
- The textbook calls the tree B-tree instead of B<sup>+</sup>-tree. In some other textbooks, B-tree refers to the variant where the actual records are kept at internal nodes as well as the leaves. Such a scheme is not practical. Keeping actual records at the internal nodes will limit the number of keys stored there, and thus increasing the number of tree levels.

# Searching Example

- Suppose that we want to search for the key K. The path traversed is shown in bold



# Searching Algorithm




- Let  $x$  be the input search key.
- Start the searching at the root
- If we encounter an internal node  $v$ , search (linear search or binary search) for  $x$  among the keys stored at  $v$ 
  - If  $x < K_{\min}$  at  $v$ , follow the left child pointer of  $K_{\min}$
  - If  $K_i \leq x < K_{i+1}$  for two consecutive keys  $K_i$  and  $K_{i+1}$  at  $v$ , follow the left child pointer of  $K_{i+1}$  (same as the right child pointer of  $K_i$ )
  - If  $x \geq K_{\max}$  at  $v$ , follow the right child pointer of  $K_{\max}$
- If we encounter a leaf  $v$ , we search (linear search or binary search) for  $x$  among the keys stored at  $v$ . If found, we return the entire record; otherwise, report not found.

# Insertion Procedure


- Suppose that we want to insert a key  $K$  and its associated record.
- Search for the key  $K$  using the search procedure
- This will bring us to a leaf  $x$
- Insert  $K$  into  $x$ 
  - Splitting (instead of rotations in AVL trees) of nodes is used to maintain properties of B<sup>+</sup>-trees [next slide]

# Insertion into a Leaf

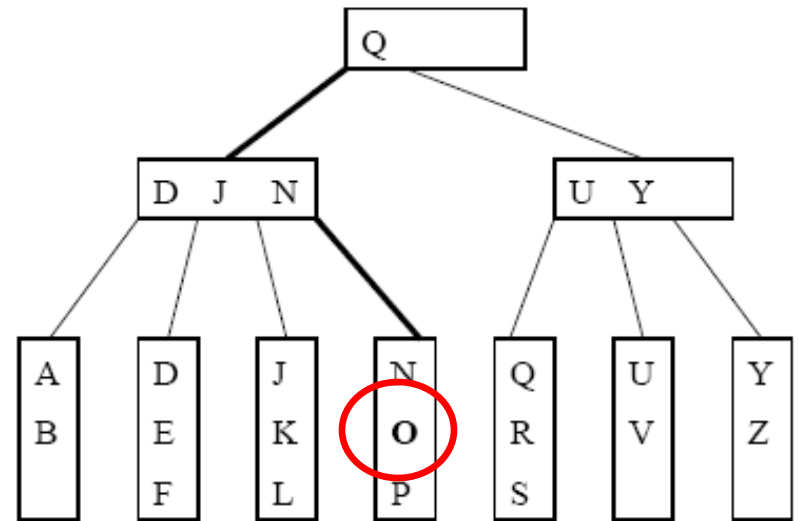


- If **leaf**  $x$  contains  $< L$  **keys**, then insert  $K$  into  $x$  (at the correct position in node  $x$ )
- If  $x$  is **already full** (i.e. containing  $L$  keys) then **split**  $x$ 
  - Cut  $x$  off from its parent
  - Insert  $K$  into  $x$ , pretending  $x$  has space for  $K$ . Now  $x$  has  $L+1$  keys.
  - After inserting  $K$ , split  $x$  into 2 new leaves  $x_L$  and  $x_R$ , with  $x_L$  containing the  $\lfloor (L+1)/2 \rfloor$  **smallest keys**, and  $x_R$  containing the **remaining**  $\lceil (L+1)/2 \rceil$  **keys**. Let  $J$  be the minimum key in  $x_R$
  - Make a copy of  $J$  to be the parent of  $x_L$  and  $x_R$ , and insert the copy together with its child pointers into the old parent of  $x$ .

# Inserting into a Non-full Leaf (L=3)




Search for O.




Insert O and maintain the order.



# Splitting a Leaf: Inserting T




Search for T.




Insert T and the B+ tree condition is violated.

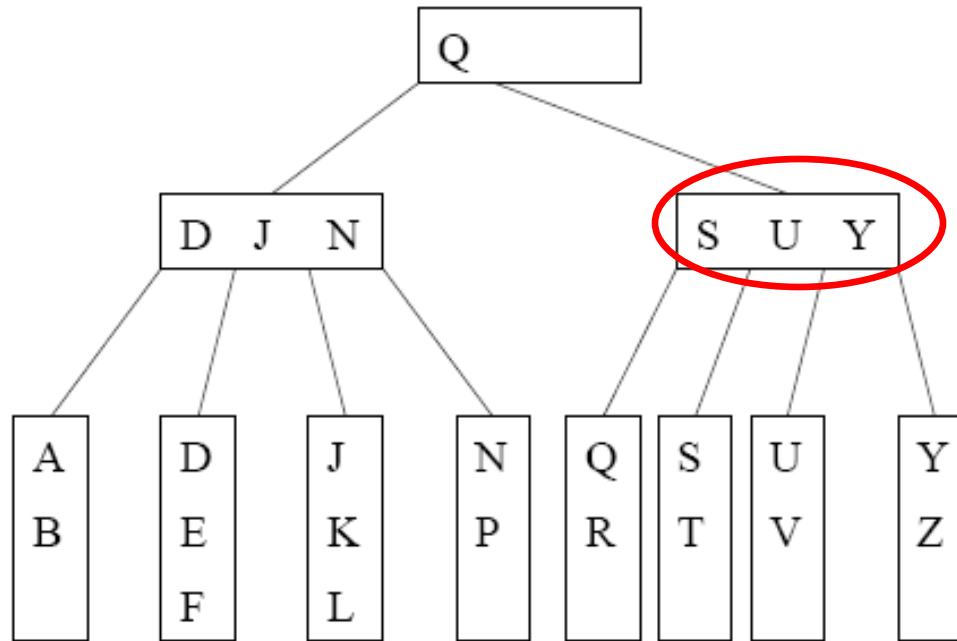
# Splitting Example 1



Split the leaf and distribute the keys.



Make S the parent of the two new leaves.




Insert S into the parent.


Maintain the order of keys and child pointers.

- Two disk accesses to write the two leaves, one disk access to update the parent
- For  $L=32$ , two leaves with 16 and 17 items are created. We can perform 15 more insertions without another split

# Splitting Example 2




Search for M.




Insert M and the B+-tree condition is violated.

# Cont'd



Split the leaf and distribute the keys.



Make L the parent of the two new leaves.

However, we cannot just insert L into the parent as it is already full.


**=> Need to split the internal node**

# Splitting an Internal Node


To insert a **key K** into a full **internal node x**:

- Cut x off from its parent
- Insert K and its left and right child pointers into x, pretending there is space. Now x has M keys (and M+1 pointers).
- **Split x** into 2 new internal nodes  $x_L$  and  $x_R$ , with  $x_L$  containing the  $(\lceil M/2 \rceil - 1)$  **smallest keys**, and  $x_R$  containing the  $\lfloor M/2 \rfloor$  **largest keys**. Note that the  $(\lceil M/2 \rceil)^{\text{th}}$  key J is not placed in  $x_L$  or  $x_R$
- Make J the parent of  $x_L$  and  $x_R$ , and insert J together with its child pointers into the old parent of x.

# Example: Splitting Internal Node (M=4)



Insert L and its child pointers into the parent.




Split the parent.

The key J becomes the parent of the two new internal nodes.

Insert J into the next parent.

# Cont'd



Split the parent.

The key J becomes the parent of the two new internal nodes.

Insert J into the next parent.



# Termination

- Splitting will continue as long as we encounter full internal nodes
- If the split internal node  $x$  does not have a parent (i.e.  $x$  is a root), then create a new root containing the key  $J$  and its two children

